

# No Subject

From: **no-reply@qut.edu.au** | no-reply@qut.edu.au

Wednesday, 26 Oct at 3:26 PM

To: **Alex McDermott** | alex.mcdermott@connect.qut.edu.au



Hi Alex,

Thank you for your assignment extension request (**FORM-AEX-275605**).

We have approved your request and the due date for your assignment **Progress Report**, for unit EGH400-1 has been extended by 48 hours from the original due date. If your unit outline does not specify that your assignment is eligible for an extension, this confirmation email is not valid and unless you submit by the original due date, the late assessment policy will apply.

You are responsible for ensuring that this assignment is eligible for extension before submitting it after the original due date. Check your [unit outline](#) for eligibility.

Be aware that a copy of this email is kept on file. You should not alter this email in any way. Email notifications that have been altered or differ in any way from the original may result in an allegation of student misconduct as set out in the [Student Code of Conduct](#).

**Need extra support?** You can access free, confidential [counselling with qualified professionals](#). We also offer [planning and support if you have a disability, injury or health condition](#) that affects your ability to study. If you are struggling to meet your university commitments, [academic support](#) is also available.

**Have a question?** You can contact us by email or phone. We're also available to help you on campus or via online chat. Visit the [HiQ website](#) for our contact details and opening hours.

 **Email us**

 **+61 7 3138 2000**

HiQ is your place to go for **student services, support** and **general enquiries**: [qut.to/abouthiq](https://qut.to/abouthiq)

hi, how can we help you?

You have received this email because you have submitted an assignment extension request. View our [Privacy and Security statement](#).

Ref ID: 10494367 FORM-AEX-275605

# Creating a Voxel Rendering Engine

EGH400-1 Progress Report - Semester 2, 2022

Supervised by Professor Clinton Fookes

**Alex McDermott (N10494367)**

Queensland University of Technology  
October 28, 2022

## Contents

<b>1</b>	<b>Current Circumstances</b>	<b>2</b>
<b>2</b>	<b>Project Summary</b>	<b>2</b>
<b>3</b>	<b>Interim Deliverables</b>	<b>2</b>
3.1	Language Choice . . . . .	2
3.2	Wgpu . . . . .	2
3.3	Ray Marching . . . . .	3
3.4	Bevy . . . . .	5
<b>4</b>	<b>Risk Management</b>	<b>6</b>
<b>5</b>	<b>Ethical Considerations</b>	<b>6</b>
<b>6</b>	<b>Sustainable Development Principles</b>	<b>7</b>
<b>7</b>	<b>Key Stakeholders</b>	<b>7</b>
<b>A</b>	<b>Deliverables and Timeline</b>	<b>7</b>

## List of Figures

1	Definition of vertices and colours written to the GPU buffer. . . . .	3
2	Single triangle coloured using vertex colours. . . . .	3
3	Illustration demonstrating the concept of ray marching. . . . .	4
4	Ray marched sphere using Rust and Wgpu. . . . .	5
5	Basic 3D scene created with Bevy. . . . .	5
6	Ray marching shader on a cube, giving the illusion of a volume. . . . .	6

## Acronyms

<b>API</b>	Application Program Interface. <a href="#">2</a> , <a href="#">3</a>
<b>CPU</b>	Central Processing Unit. <a href="#">2</a>
<b>GPU</b>	Graphics Processing Unit. <a href="#">1</a> , <a href="#">2</a> , <a href="#">3</a> , <a href="#">5</a> , <a href="#">7</a>
<b>QUT</b>	Queensland University of Technology. <a href="#">2</a> , <a href="#">7</a>
<b>SDF</b>	Signed Distance Function. <a href="#">4</a> , <a href="#">8</a>
<b>WGSL</b>	WebGPU Shading Language. <a href="#">3</a> , <a href="#">6</a>

# 1 Current Circumstances

I'd like to preface this report with a short summary of the current circumstances. This topic, being the creation of a voxel engine, has only been the topic for of this thesis for just under three weeks. Up until Week 10, another student and myself where working on a different topic which was supposed to be through an industry connection with Airbus. Sadly, they where unable to follow through on their promise to find a representative to manage the project on their end so this arrangement fell through. Fortunately, our [Queensland University of Technology \(QUT\)](#) supervisor Professor Clinton Fookes, has been very accommodating towards this jarring series of events so late in the semester. As a result of this, my current project is a student led idea that I pitched as an alternative that will be more feasible due to the much reduced time frame at hand. I chose this topic has I have past experience with similar rendering techniques such as Ray Marching and feel that this somewhat familiar field and interest will aid in making up for lost time.

## 2 Project Summary

The goal of this thesis project is to create voxel rendering engine. Traditionally, rendering has been done with the use of triangles. These triangles are connected together in varying densities and sizes to create 3D models. They are then projected from 3D space to the 2D camera plane, and through a process called rasterisation the pixels they occupy are shaded. To begin understanding how voxel rendering differs from this we must first define what a voxel is. A voxel is a cell of a 3D grid where each cell stores information about the density at that particular location in space. Essentially a 3D pixel. Due to its innate ability to store this density data it is commonly used in medical imaging as an output data structure for Magnetic Resonance Imaging and Computed Tomography scans. Another use for voxels is in video games where they can be be used to implement destructible terrain such as in Teardown [1]. The aim for this thesis is to create a program that is able to render these objects interactively, in real time with lighting and shading effects.

## 3 Interim Deliverables

### 3.1 Language Choice

Rendering is generally a very computationally expensive process no matter the algorithm that is being used. This is due to the sheer amount of data that needs to be processed/created in order to produce an image of a usable size. Considering most displays these days output images at a minimum of 1080p meaning 1920 pixels wide by 1080 pixels high. For a renderer creating one of these images, this alone is already two million pixel colour calculations that is likely being performed at least 30 times per second in a real time application. Observing the need for performance in creating a real time rendering application it was clear an efficient low level language would be needed. Due to this consideration I chose Rust to make this application. It's a modern, fast, memory safe, low level language with performance comparable to C although with a significantly more pleasant developer experience. Rust ranked number one on the list of most loved languages in the most recent StackOverflow Developer Survey [2] and with such an active community, has a wide array of high quality packages readily available. Also, due to this languages overwhelming popularity, there is a wealth of educational resources available to aid in development.

### 3.2 Wgpu

Despite having chosen a sufficiently fast programming language above, implementing a real time rendering engine that runs entirely on the [Central Processing Unit \(CPU\)](#) is just not feasible. To accomplish this project we will be taking advantage of the [Graphics Processing Unit \(GPU\)](#) which is able to process many pixels in parallel, thus improving computation times per frame by orders of magnitude. Although, there are many different [GPU](#) hardware architectures and [Application Program Interface \(API\)](#) options making it impractical to target all of them as a single developer. Because of this, it is common to use an abstraction layer such as Wgpu [3], which I have opted to use here, that handles the communication between these platforms under the hood. Wgpu is a Rust library the implements the WebGPU [4] [API](#), meaning it exposes a single set of methods that give us the ability to run natively in Vulkan, Metal, DirectX, and OpenGL graphics environments.

Having decided on a language and graphics [API](#) it was time to begin getting content drawn to the screen. I began this processing by following the [Learn Wgpu \[5\]](#) tutorial. This walked through the process of setting up a Rust project with the required dependencies as well as instantiating a window for the application with a basic event loop. Given this window object, a few more components are required before we could obtain a visual output. Firstly an adapter is created which handles interactions with the graphics card and provides information such as its name, backend and power preference. The adapter can then be used to create a surface which is the part of the window that can be drawn too. Having this level of control over the graphics process means things such as resizing and updating of the window dimensions need to be kept track of manually so this was also implemented here.

With these key components created I can now explain the [GPU](#) pipeline. This is the steps the graphics system takes in order to go from a set of vertices to an image on the screen. In this simple case we are just looking to get up and running so I started with a single triangle. A buffer, a contiguous location of memory located on the [GPU](#), was created to hold the required vertices for our primitive shape. This can be seen in the figure below along with the colours assigned to each vertex.

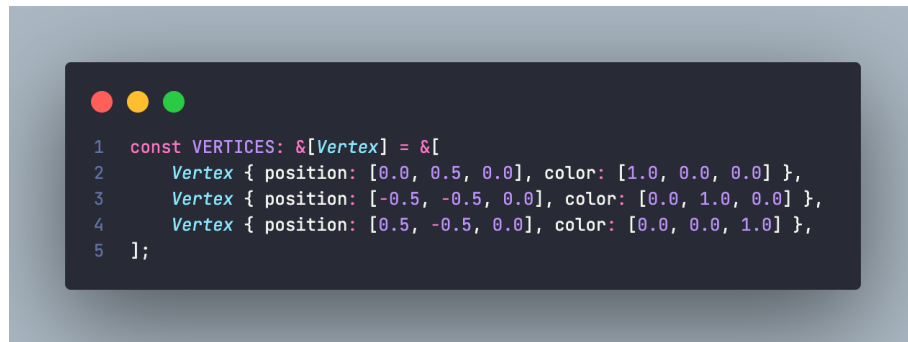


Figure 1: Definition of vertices and colours written to the [GPU](#) buffer.

This data is then passed into our pipeline, which in the case of Wgpu is handled by the [WebGPU Shading Language \(WGSL\)](#). Due to the simplistic nature of our shader, no processing is required on this vertex data and since the colours are already included, they can be extracted and used in the fragment shader with a small amount of code. This results in the image seen in the figure below.



Figure 2: Single triangle coloured using vertex colours.

### 3.3 Ray Marching

Moving on from the basic Rust [GPU](#) rendering proof of concept explored above, it was time to get familiar with the semantics of [WGSL](#). Having experience with the older OpenGL Shading Language, I was interested in figuring out the similarities and differences so I could apply my existing knowledge

to this new system. To do this, my idea was to implement a basic but familiar ray marching render. A benefit of this was that, from my understanding, it seemed like the fundamentals of ray marching were quite similar to that of voxel rendering so I was hopeful some code, as well as knowledge would be transferable. Building off the triangle that was currently being rendered, the colour data was stripped out and an extra vertex was added to create a rectangle that filled the whole screen. The reasoning for this was that ray marching is a software technique, it could be entirely implemented in the fragment shader. Each pixel of the image could be calculated separately and displayed on a single full screen rectangle.

Discussing the technical details of ray marching, the whole premise is based around a function that returns the distance to the closest object in the scene given any sample point. This function is referred to as a [Signed Distance Function \(SDF\)](#). For each pixel, a ray is sent out at an angle determined by the field of view. The [SDF](#) is then sampled to determine the minimum safe distance the ray can be marched before intersection with an object. This process is repeated until the [SDF](#) returns a value less than some threshold, indicating an object has been hit, or reaches some set number of maximum iterations indicating the opposite. This concept is visualised particularly well in the figure shown below.




Figure 3: Illustration demonstrating the concept of ray marching.

The process described here was implemented in the fragment shader of our graphics pipeline. For each pixel, if a hit was detected, the gradient of the [SDF](#) was found by calculating the differences between two samples a small distance apart. This gradient was used as the basis for Phong shading [6] which allows the sphere below to be visualised with approximated diffuse lighting and specular highlights.




Figure 4: Ray marched sphere using Rust and Wgpu.

### 3.4 Bevy

Observing the output of the above sections, the result is quite desirable. It provides us with a means of writing shader code for our GPU pipeline and visualising the output. Although, there was a number of steps involved in reaching this point, primarily all the manual interfacing with the Wgpu library which created a large amount of boilerplate code. My concern with this was that, as our application grows more complex, these more manual processes would take more and more time from actual engine development. To combat this, I decided to make the switch to a very popular Rust game engine called Bevy [7]. Bevy is a data driven game engine that uses an Entity Component System which has become a common approach in previous years. Even early on the benefits of this switch would be apparent, allowing for the remove all of the Wgpu boilerplate code. Very conveniently Bevy comes with plugins for window creation, event loops and rendering primitives out of the box. It also has a system in place for writing custom shaders for objects which we will be using extensively. This switch would allow us to focus on writing the actual shader code instead of micro optimising the GPU surfaces, buffers and etc. Further utilising the prospering Rust package ecosystem, I was able to import an open source first person flying camera controller [8]. This would prove to be very helpful as it would allow for debugging of vector maths with relation to object and camera positions. With less than 50 lines of code using the Bevy ECS system we are able to create the basic 3D scene as shown below as opposed to the hundreds required previously to draw a single quad using Wgpu.




Figure 5: Basic 3D scene created with Bevy.

Having now switched libraries, I was interested in porting over the shader code I previously wrote for

Wgpu. Fortunately both libraries are quite progressive in terms of shader language choice and supported [WGSL](#). This made bring over the bulk of the shading code quite simple as all that needed changing was how the ray directions where being calculated. Having watched a video by the Tuxedo Labs team on the technical aspects of their game Teardown [9], I saw how they actually used cubes meshes to render voxel volumes. This was also a reason behind the switch to Bevy as with the new system, it would be much easier to create many cube meshes if required in the future. This could allow for a scene being composed of multiple voxel volumes without manually specifying all the vertex buffers and etc. In terms of converting our code to render our ray marched sphere on the texture of a cube, as stated the ray direction would need to be recalculated. It was found that this could be done by accessing a points 3D world position from within the fragment shader and then subtracting the cameras position. The result of this "projection" of the ray marched sphere onto the cube gives the illusion of a volume as that is effectively what is being simulated. Shown below is the result of this process, the black cube outline can easily be hidden but is left in to enhance understanding.




Figure 6: Ray marching shader on a cube, giving the illusion of a volume.

## 4 Risk Management

Thankfully, being a software application, there is a negligible chance of physical risks associated with this project. Despite this, there is still some minimal data related risks present which have been accounted for as follows. It is widely known, but not often accounted for, that data stored in a single place can be susceptible to loss due to hardware failure or misplacement of the physical device. To combat this, Git version control has been used to both keep track of incremental program changes on a line by line basis, as well as keep a cloud based backup through the GitHub service [10]. Another crucial aspect that must be considered when working on a programming project is the data involved and the security and privacy expected when handling it. In the case of this voxel engine, there is no sensitive data being handled currently. There is a chance that this could change in the future if the application is given the ability to read in medical data. In the event that this happens I will be sure to implement it professionally and utilise security best practices.

## 5 Ethical Considerations

Ethics are an important consideration when completing any project. For Engineers, it is crucial we look to the Engineers Australia Guidelines [11] to aid in our professional conduct. Firstly, regarding codes 1.1, 1.3, 3.1, and 3.3, in summary, one must act in good conscience, respect dignity, uphold trustworthiness, and honestly communicate. This has been upheld throughout this project via diligent citation and



referencing of a vast number of helpful resources, giving credit where credit is due. The codes described under section 2 for competent practising have also been upheld by my continued effort to stay up to date with the constantly evolving field of software and technology. I do this out of passion but also in the interest of always striving to produce a clean, maintainable software solution that follows best practices.

## 6 Sustainable Development Principles

Sustainability is another key aspect that must be taken into account when undertaking any engineering project. As such, careful consideration has been taken thus far in this thesis to make conscious decisions in this regard. Being a small software based project it is fortunate that this project has a very limited impact in terms of environmental sustainability. The one main impact it does have however is the power consumption of the hardware it is run on. In an attempt to minimise this it is important that developers such as myself strive to write efficient code that effectively uses compute cycles. So far this has been implemented in this project by utilise the [GPU](#) for per pixel parallel computing.

Another important aspect of sustainability when completing an engineering project is being able to sustain motivation and development drive. Burnout is a common issue within the industry and seeing that development is conducted in a motivationally sustainable manner is crucial to its mitigation. Throughout this project this was achieve by taking manageable step sizes towards out goal. It would be unrealistic to being working with a new language, library, and concept and expect to not get overwhelmed. This is the reason I started with the familiar concept of ray marching, to get my footing with both the language and the library being used. From here I have began pivoting towards implementing the voxel features having now become acquainted with these new technologies in a more friendly environment.

The final aspect of sustainability that was considered for this project was the codes maintainability. As any developer would know, when coming back to one of your own projects, or picking up another developers project, it can take a significant amount of time to become acquainted with the code base and understand what is going on. This is doubly true if the code was poorly write and documented. Taking note of this, in order to support sustainable development and avoid large refactoring and rewrites, it was a goal to write clean and sustainable code from the beginning. This means comments within the source code, meaningful commit messages in version control, progress documentation such as this, as well as clean code structure with sensible variable names.

## 7 Key Stakeholders

After accounting for the issues we had with our industry partners, there is now fewer stakeholders in this project. Aside from myself and the academic experience this project will impart on me, there are no other third parties that are significantly invested in the outcome of this thesis. My supervisor, Professor Clinton Fookes, would hold a limited amount of stake in this project as its success or failure could potential reflect on his supervisory presence. Although his support so far through this projects late changes has been warmly welcomed. Since this has ultimately changed to be a student led project, I suspect his and the investment of [QUT](#) will ultimately be limited.

## A Deliverables and Timeline

The table included below outlines the deliverables that have currently been completed as part of this project as well as those that are yet to be delivered. It also specifies their dependants along with expected completion date. A number of these items are outlined by [QUT](#) as part of completing an honours thesis in engineering and help provide good structure to the research method being followed. Moving to the interim milestones, those marked as completed have been detailed as part of this project report whereas the others will be explored here.

Firstly, upon coming back to this project in semester 1 of next year, I am to begin by generating a sphere made of voxels. This should be an attainable goal and will allow me to get to grips with the voxel rendering technique without having to learn a volume file format at the same time. Once I have this sphere being created, will need to update the current ray marching render to traverse this volume instead

of the SDF. This will require modification of the algorithm as well as primarily the shading code in order to display the generated volume on the screen. Following this, and likely the final interim deliverable for this thesis project, I am to add the ability to load in external models. This could open up the engine for many users, allowing for visualisation of voxel game assets, medical scans and more.

Number	Name	Due	Dependency	Type	Status
1	Project Proposal	W07 Sem 2, 2022	N/A	Assessed	Completed
2	Get shader working	W10 Sem 2, 2022	1	Interim	Completed
3	Implement ray marching	W11 Sem 2, 2022	2	Interim	Completed
4	Render volume to cube	W12 Sem 2, 2022	3	Interim	Completed
5	Progress Report	W13 Sem 2, 2022	4	Assessed	Completed
6	Oral Presentation	W03 Sem 1, 2022	5	Assessed	Incomplete
7	Generate voxel sphere	W06 Sem 1, 2023	6	Interim	Incomplete
8	Update shading for voxels	W09 Sem 1, 2023	7	Interim	Incomplete
9	Load external voxel volume	W12 Sem 1, 2023	8	Interim	Incomplete
10	Final Project Report	W13 Sem 1, 2023	9	Assessed	Incomplete
11	Final Oral Presentation	End Sem 1, 2023	10	Assessed	Incomplete

## References

- [1] T. Labs, “Teardown steam,” 2021. [Online]. Available: <https://store.steampowered.com/app/1167630/Teardown>
- [2] StackOverflow, “Stackoverflow developer survey,” 2021. [Online]. Available: <https://insights.stackoverflow.com/survey/2021>
- [3] Wgpu, “Wgpu source code,” 2022. [Online]. Available: <https://github.com/gfx-rs/wgpu>
- [4] W3C, “Webgpu specification,” 2022. [Online]. Available: <https://www.w3.org/TR/webgpu>
- [5] Wgpu, “Learn wgpu,” 2022. [Online]. Available: <https://sotrh.github.io/learn-wgpu>
- [6] Scratchapixel, “The phong model, introduction to the concepts of shader, reflection models and brdf,” 2022. [Online]. Available: <https://www.scratchapixel.com/lessons/3d-basic-rendering/phong-shader-BRDF>
- [7] Bevy, “Bevy,” 2022. [Online]. Available: <https://bevyengine.org>
- [8] sburris0, “Bevy flycam,” 2022. [Online]. Available: [https://github.com/sburris0/bevy\\_flycam](https://github.com/sburris0/bevy_flycam)
- [9] T. Labs, “Tuxedo labs teardown technical teardown,” 2020. [Online]. Available: <https://www.youtube.com/watch?v=0VzE8ROwC58>
- [10] GitHub, “Github,” 2022. [Online]. Available: <https://github.com>
- [11] E. Australia, “Code of ethics,” 2022. [Online]. Available: <https://www.engineersaustralia.org.au/publications/code-ethics>